ANALYTICS EXPERIENCE 2016

Getting Started with SAS and Hadoop

Jeff Bailey

#analyticsx

Why Hadoop?

HOW MUCH DOES THIS DRIVE COST?

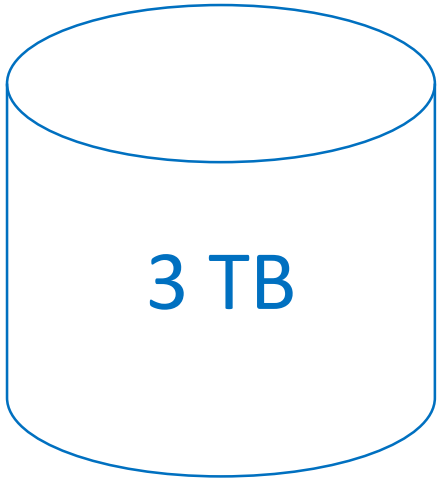

HOW MUCH DOES THIS DRIVE COST?

Silly, you couldn't get a 3TB drive in 1980!

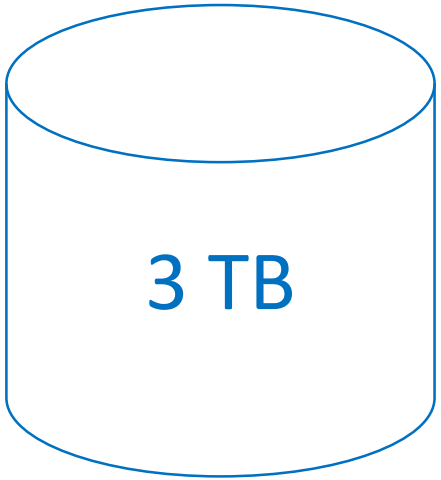

1980

$1,312,500,000

HOW MUCH DOES THIS DRIVE COST?

That's $0.03 per GB!

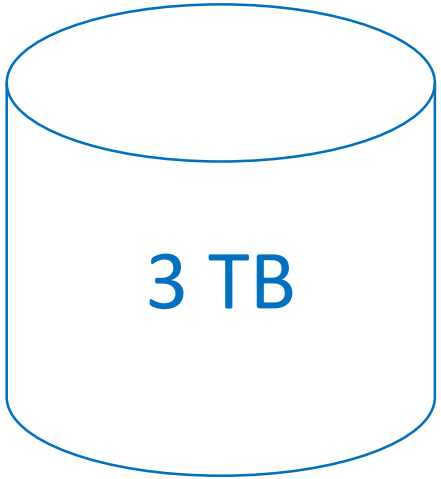

TODAY	$69
2010	$270
2005	$3,720
2000	$33,000
1995	$3,360,000
1990	$33,600,000
1985	$315,000,000
1980	$1,312,500,000

HOW MUCH DOES THIS DRIVE COST?

That's $0.03 per GB!

TODAY	$69
2010	$270
2005	$3,720
1980	$1,312,500,000

Insight: Disk Space is FREE!

IT'S NOT JUST ABOUT COST!

How long does it take to read **3 TB of data**?

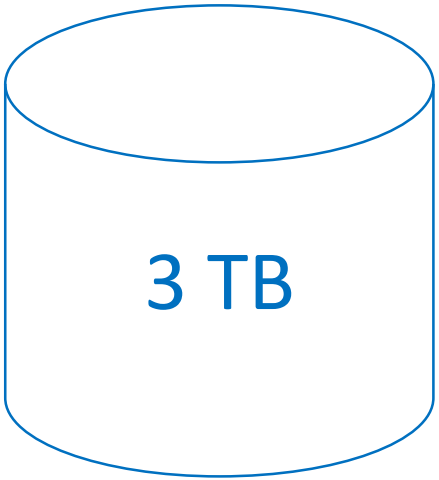

IT'S NOT JUST ABOUT COST!

How long does it take to read **3 TB of data**?

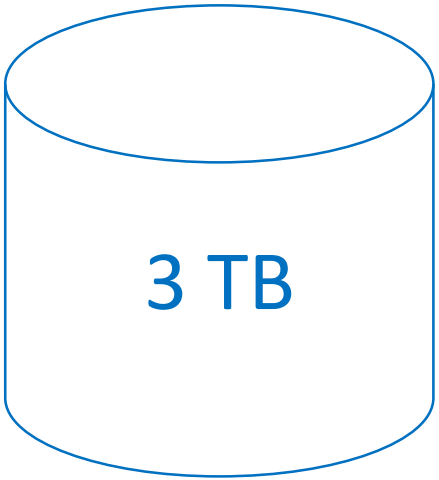

4.17 Hours

IT'S NOT JUST ABOUT COST!

How long does it take

What happens if you add more disks?

HOW LONG DOES IT TAKE TO READ A 3 TB FILE?

HOW LONG DOES IT TAKE TO READ A 3 TB FILE?

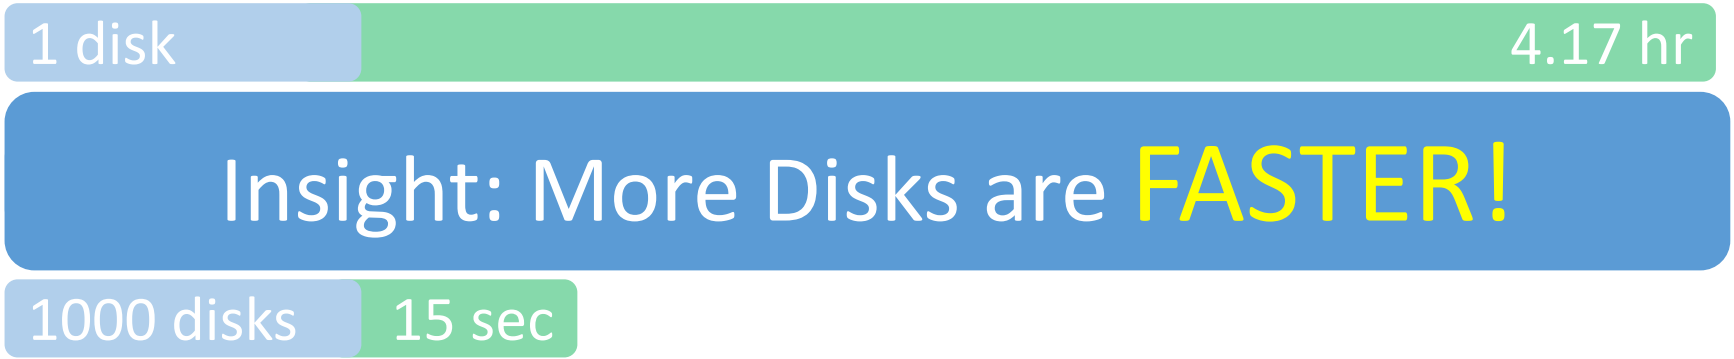

What is Hadoop?

Hadoop is a Storage Platform

- Distributed Storage Performs Great
- Data is Replicated
- Reasonable Cost
- Sits on the OS File System

Hadoop is a Processing Platform

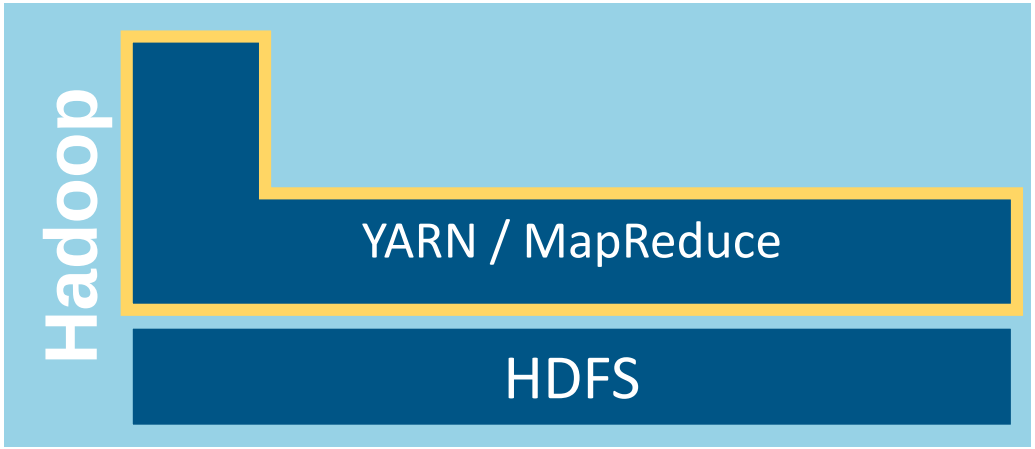

- MapReduce/YARN
- Distributed Processing
- Data Locality
- Usually Java

Apache Pig

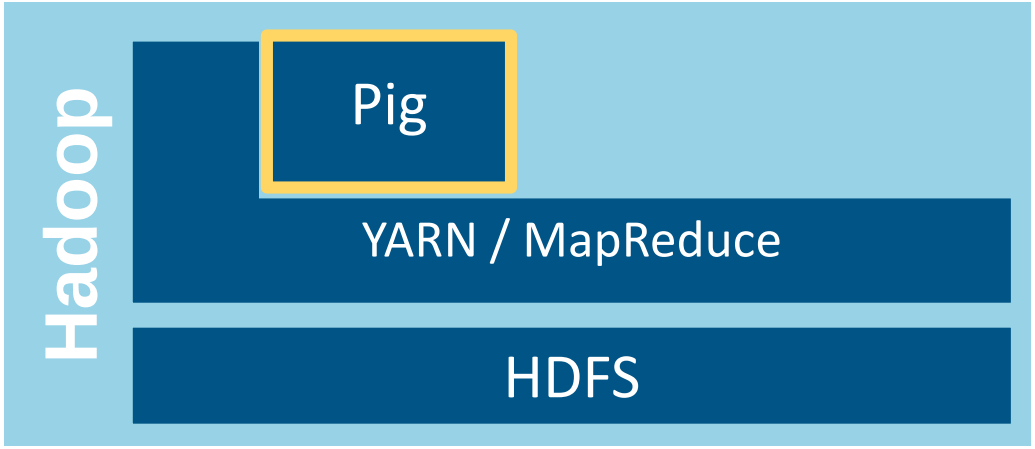

- Scripting Language
- Higher level than programming Java MapReduce
- Pig Latin scripts are converted to MapReduce jobs
- Great for joining data
- Great for transforming data

Apache Pig: Example Program

- Distributed Processing

```
people = LOAD '/user/training/customers' AS (cust_id, name);
orders = LOAD '/user/training/orders' AS (ord_id, cust_id, cost);
groups = GROUP orders BY cust_id;
totals = FOREACH groups GENERATE group, SUM(orders.cost) AS t;
result = JOIN totals BY group, people BY cust_id;
DUMP result;
```

Apache Hive

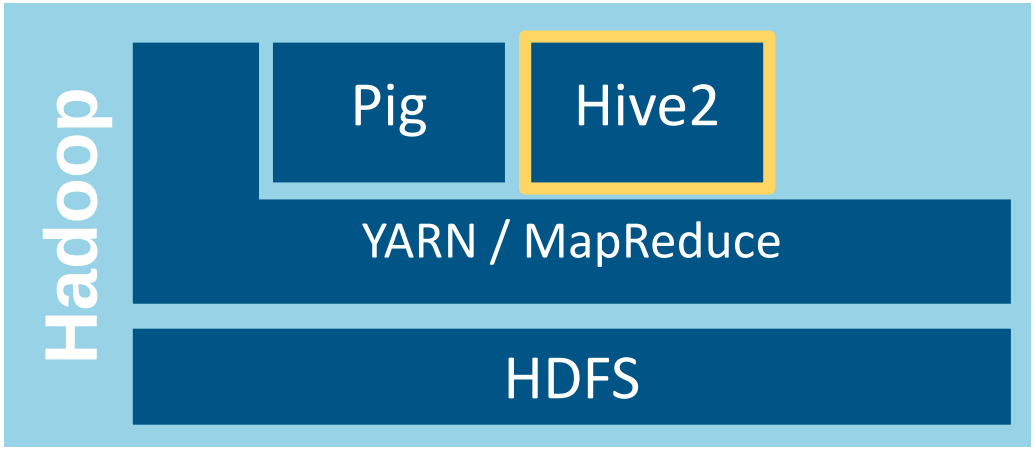

- SQL on Hadoop
- Similar to traditional SQL
- Reduces development time
- Enables BI on Hadoop
- Schema-on-Read
- You choose underlying file format

Apache Hive

- SQL on Hadoop

Hadoop

```
SELECT zipcode, SUM(cost) AS total
FROM customers
JOIN orders
ON (customers.cust_id = orders.cust_id)
WHERE zipcode LIKE '63%'
GROUP BY zipcode
ORDER BY total DESC;
```

Apache Impala is a SQL Engine

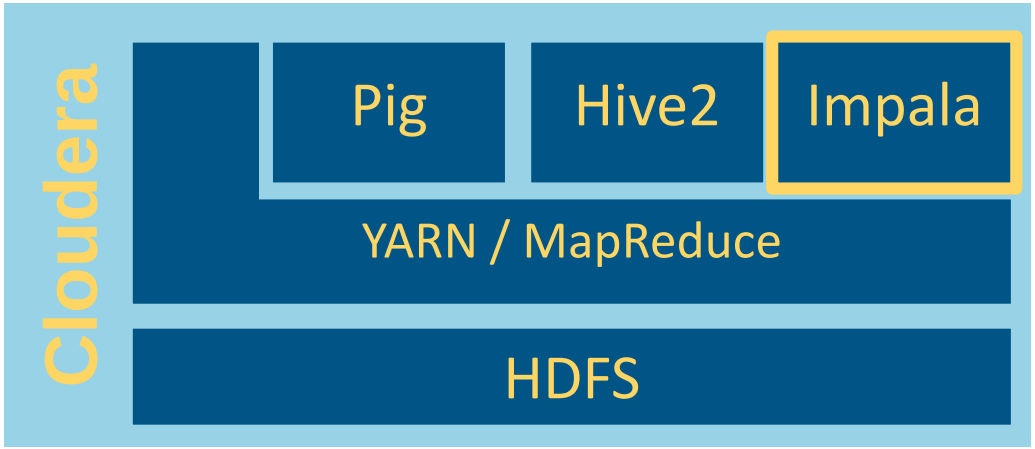

- High-performance SQL engine
- Handles concurrency well
- Does not rely on MapReduce
- Supports a dialect of SQL very similar to Hive's
- 100% open source
- Apache License

How can SAS Interact with Hadoop?

Using Base SAS 9.4 with Hadoop

#1

SAS FILENAME Statement for Hadoop

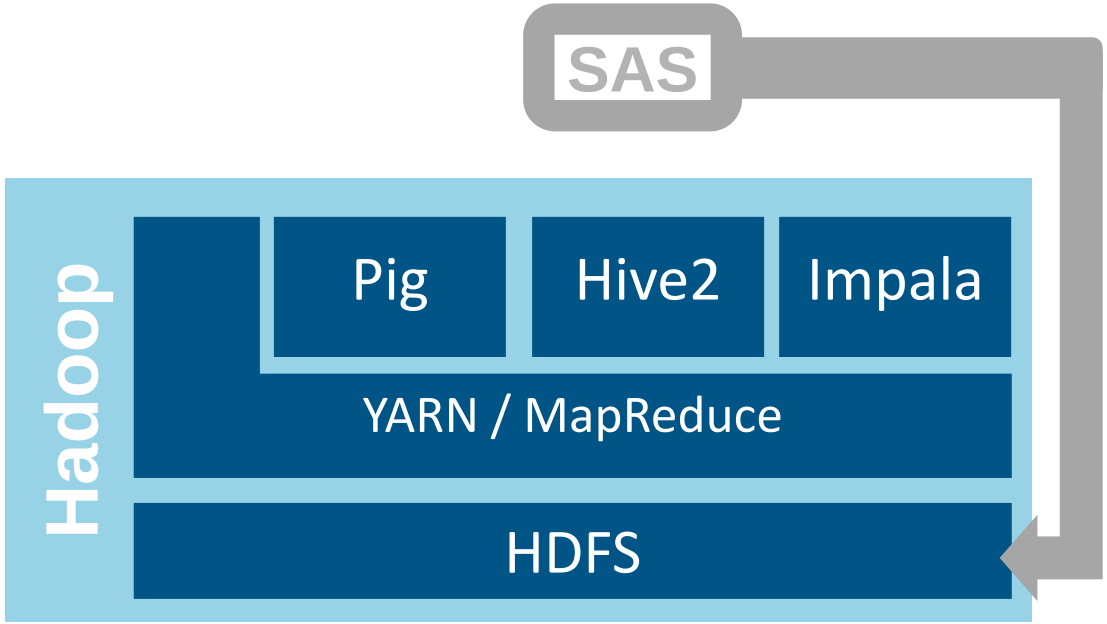

SAS FILENAME Statement for Hadoop

```
options set=SAS_HADOOP_CONFIG_PATH="\\sashq\cdh45p1";
options set=SAS_HADOOP_JAR_PATH="\\sashq\cdh45";

FILENAME hdp1 hadoop 'test.txt';

/* Write file to HDFS */
data _null_;
   file hdp1;
   put ' Test Test Test';
run;

/* Read file from HDFS */
data test;
   infile hdp1;
   input textline $15.;
run;
```

Using Base SAS 9.4 with Hadoop

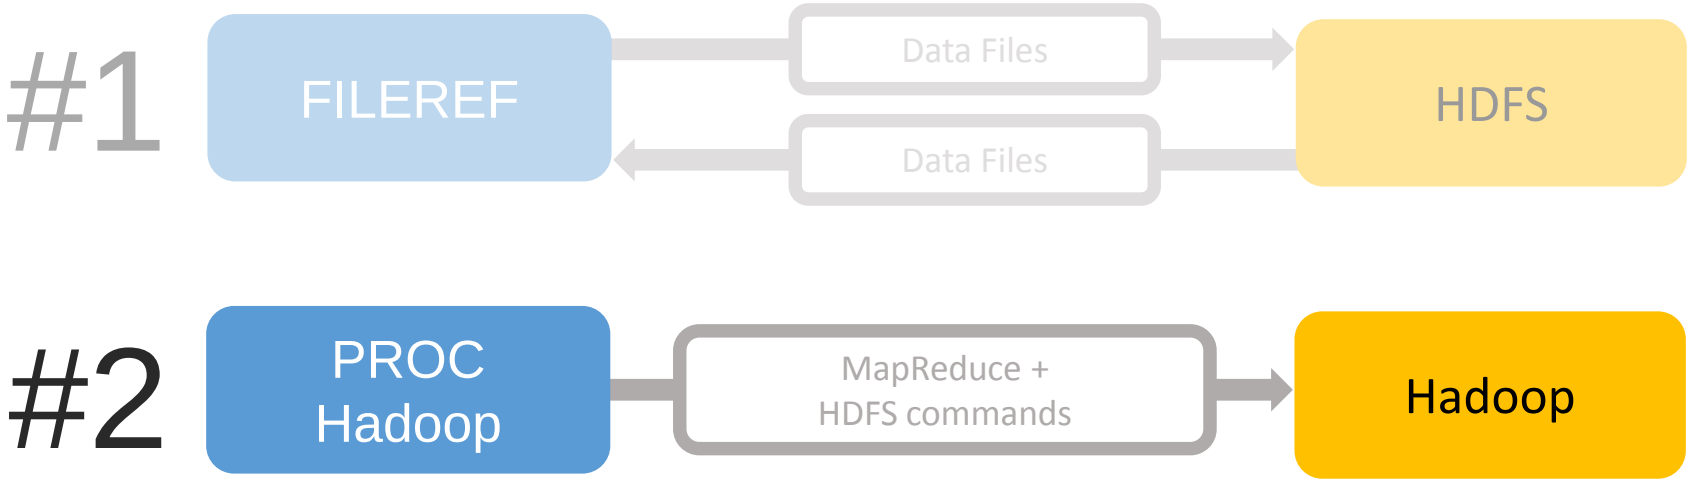

Hadoop Procedure

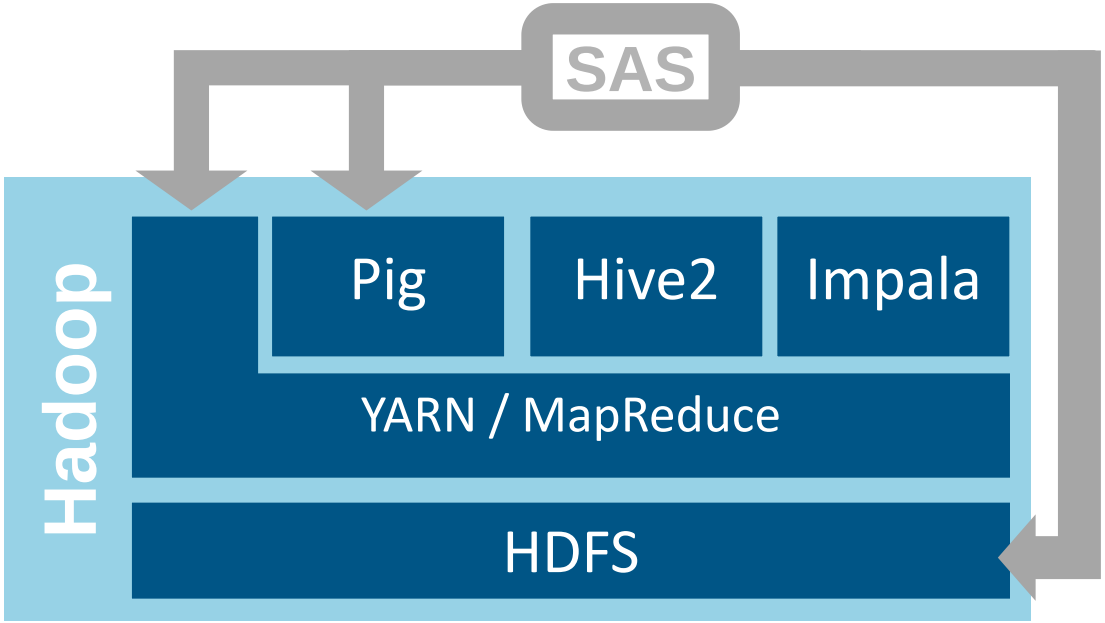

- Submit HDFS commands
- Submit MapReduce Jobs
- Submit Pig Latin programs

How Do I Submit HDFS Commands?

```
filename cfg 'C:\Hadoop_cfg\cdh57.xml';

/* Copy war_and_peace.txt to HDFS. */
/* Copy moby_dick.txt     to HDFS. */
proc hadoop options=cfg username="sasxjb" verbose;
   HDFS mkdir='/user/sasxjb/Books';
   HDFS COPYFROMLOCAL="C:\Hadoop_data\moby_dick.txt"
            OUT='/user/sasxjb/Books/moby_dick.txt';
   HDFS COPYFROMLOCAL="C:\Hadoop_data\war_and_peace.txt"
            OUT='/user/sasxjb/Books/war_and_peace.txt';
run;
```

How Do I Submit MapReduce Jobs?

```
filename cfg 'C:\Hadoop_cfg\cdh57.xml';

proc hadoop options=cfg user="sasxjb" verbose;
  mapreduce input='/user/sasxjb/Books/moby_dick.txt'
    output='/user/sasxjb/outBook'
    jar='C:\Hadoop_examples\hadoop-examples-1.2.0.1.3-96.jar'
    outputkey="org.apache.hadoop.io.Text"
    outputvalue="org.apache.hadoop.io.IntWritable"
    reduce="org.apache.hadoop.examples.WordCount$IntSumReducer"
    combine="org.apache.hadoop.examples.WordCount$IntSumReducer"
    map="org.apache.hadoop.examples.WordCount$TokenizerMapper";
run;
```

How Do I Submit Pig Latin Programs?

```
filename cfg 'C:\Hadoop_cfg\cdh57.xml';

proc hadoop options=cfg username="sasxjb‟ verbose;
     pig code=pigcode ;
run;
```

Using Base SAS 9.4 with Hadoop

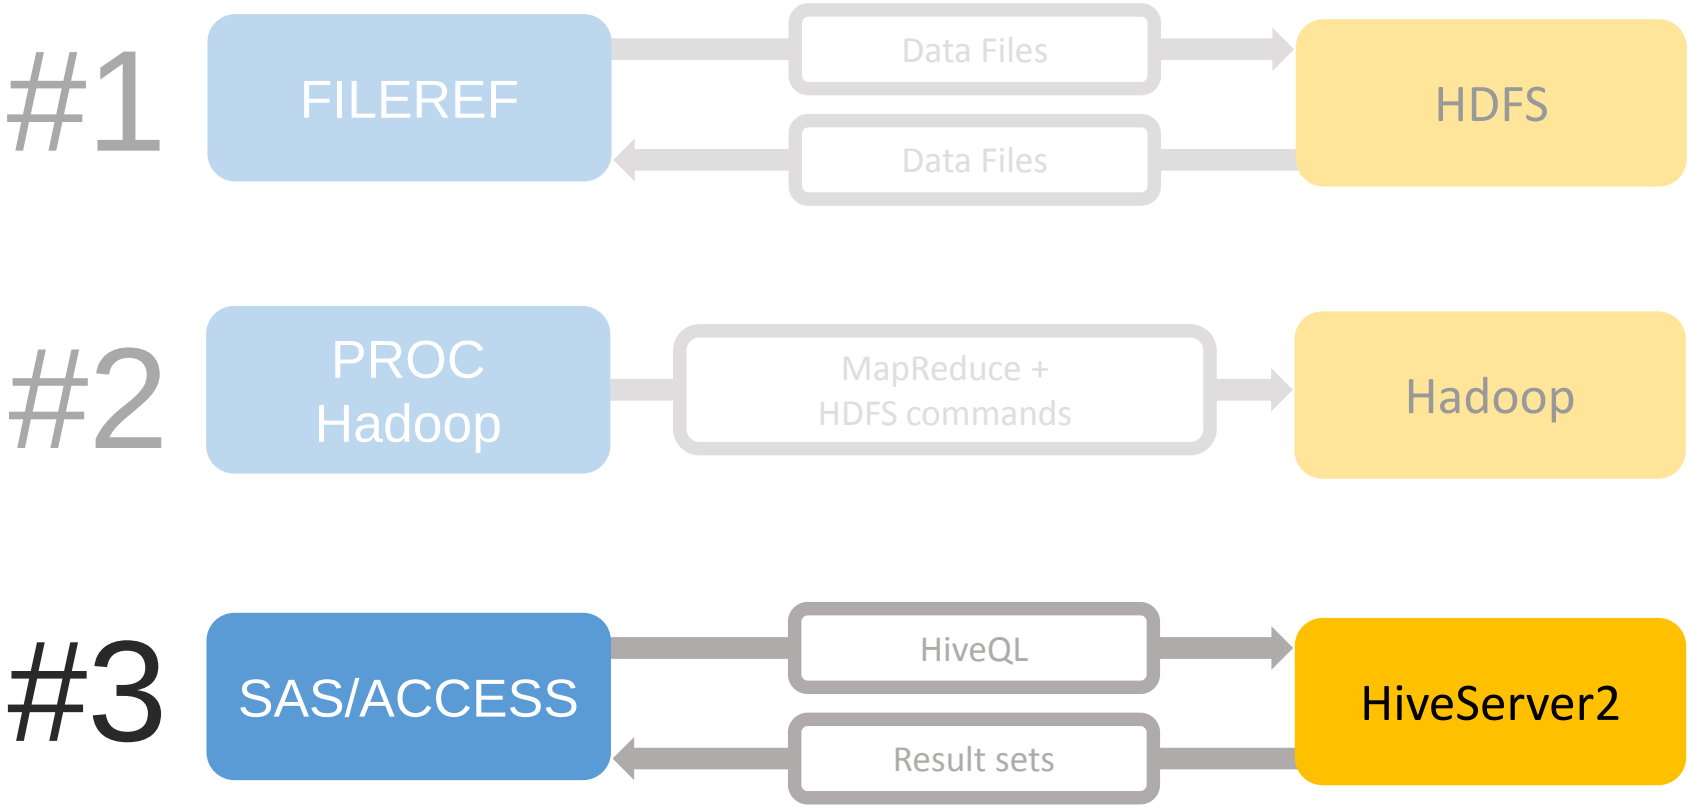

SAS/ACCESS Interface to Hadoop

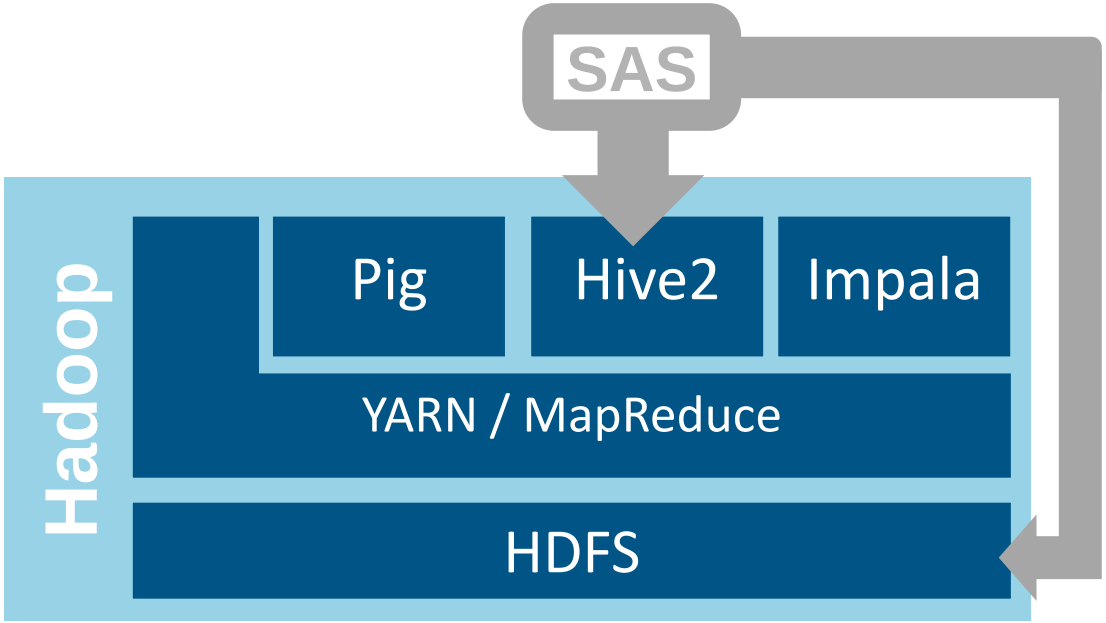

- Connects via JDBC
- Makes Hive tables look like SAS data sets
- Bulk loads directly to HDFS
- Can read directly from HDFS

How Does SAS/ACCESS Talk to Hadoop?

```
proc sql;
   select count(*) from mycdh.customer_dim
      where loyalty_program='Chocolate Club';
run;
```

?

How Does SAS/ACCESS Talk to Hadoop?

```
proc sql;
   select count(*) from mycdh.customer_dim
      where loyalty_program='Chocolate Club';
run;
```

```
select COUNT(*) from `CUSTOMER_DIM` TXT_1
WHERE TXT_1.`loyalty_program` = 'Chocolate Club'
```

SAS Generated This SQL

How Does SAS/ACCESS Talk to Hadoop?

```
proc sql;
   select count(*) from mycdh.customer_dim
      where loyalty_program='Chocolate Club';
run;
```

```
OPTIONS  SASTRACE=',,,d' SASTRACELOC=SASLOG NOSTSUFFIX;
```

```
select COUNT(*) from `CUSTOMER_DIM` TXT_1
WHERE TXT_1.`loyalty_program` = 'Chocolate Club'
```

SAS Generated This SQL

SAS/ACCESS Interface to Hadoop

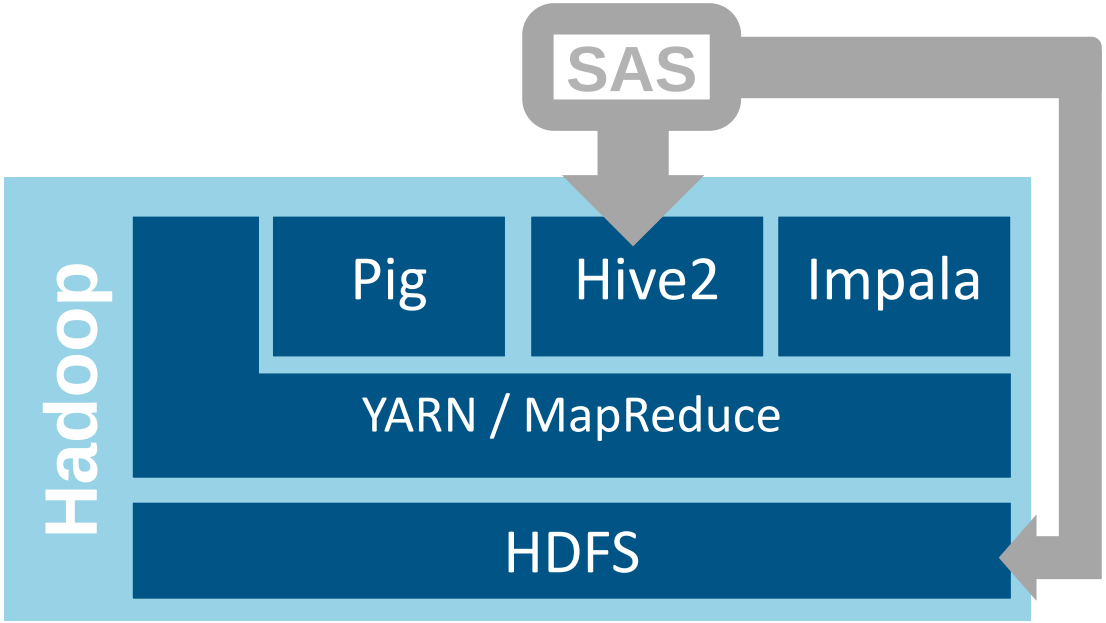

- **Generates HiveQL**
- Connects via JDBC
- Makes Hive tables look like SAS data sets
- Bulk loads directly to HDFS
- Can read directly from HDFS

We Can Write Our Own HiveQL!

```
proc sql;
   connect to hadoop (server=quickstart
                      user=cloudera);
   execute (create table store_cnt
              row format delimited
              fields terminated by '\001'
              stored as parquet
             as
                select customer_rk, count(*) as tot
                  from order_fact
                 group by customer_rk) by hadoop;
quit;
```

Explicit Pass-Through

What about Apache Impala?

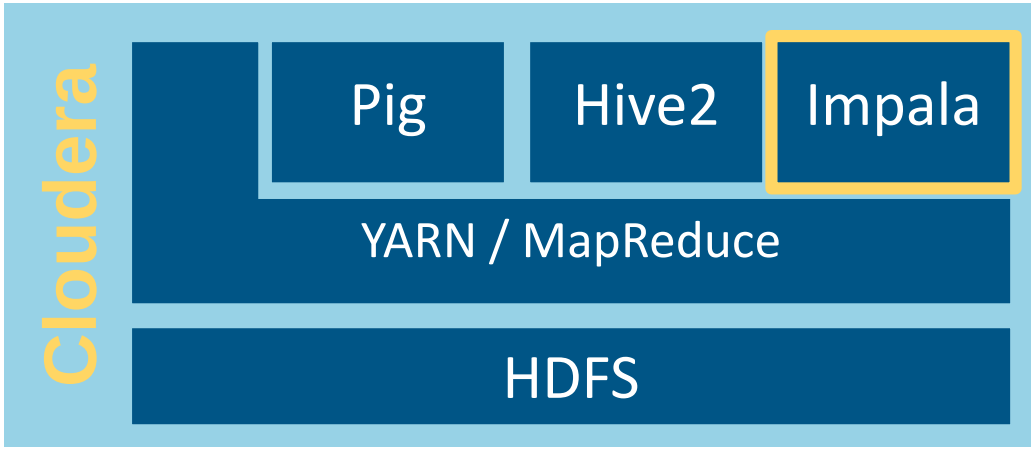

SAS/ACCESS Interface to Impala:

- Connects via ODBC
- Makes Hive tables look like SAS data sets
- Bulk loads directly to HDFS

#4

In-Database: Code Accelerator

What is SAS In-Database Code Accelerator?

```
proc ds2 indb=yes;
   thread tpgm / overwrite=yes;
      method run();
        set hdplib.intable;
        output;
   end;
   endthread;
   run;
   data hdplib.outdata

(overwrite=yes);
      dcl thread tpgm hdpdata;
      method run();
         set from hdpdata;
      end;
   enddata;
   run;
quit;
```

SAS In-Database Code Accelerators let you run SAS code inside Hadoop. With this you get:

- DS2 processing (modern DATA Step)
- More Data Types
- Code Packages
- More Programming Structures
- Parallel Database Operations
- Thread Programs Run Inside Database

In-Database Code Accelerator Runs in Hadoop

```
proc ds2 indb=yes;
   thread tpgm / overwrite=yes;
      method run();
        set hdplib.intable;
        output;
   end;
   endthread;
   run;
   data hdplib.outdata

(overwrite=yes);
        dcl thread tpgm hdpdata;
        method run();
            set from hdpdata;
        end;
    enddata;
    run;
quit;
```

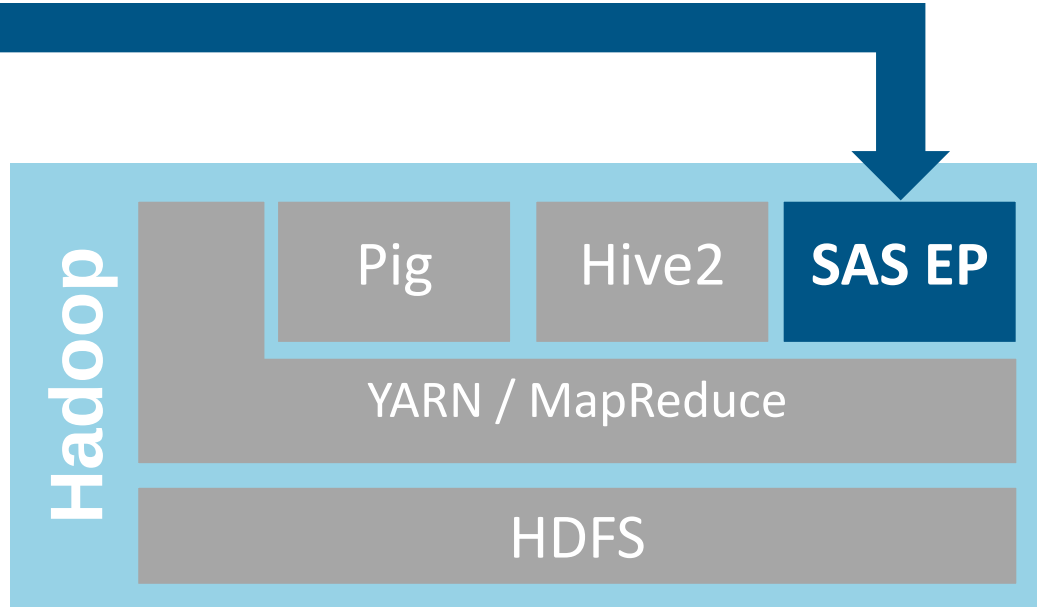

In-Database: Scoring Accelerator

What is SAS In-Database Scoring Accelerator?

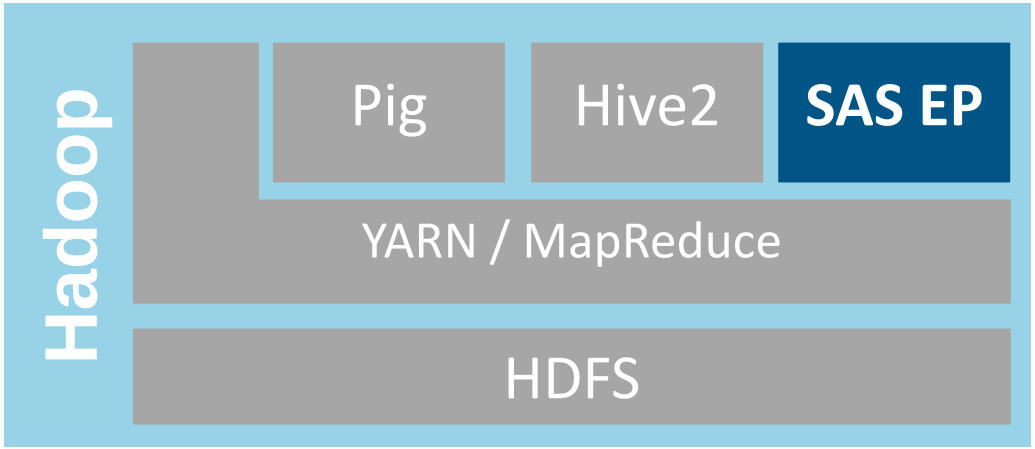

SAS In-Database Scoring Accelerator lets you score models inside the cluster. With this you get:

- Uses the SAS Embedded Process
- Faster Scoring
- Less data movement – score data where it lives
- Uses fewer resources

What does the Scoring Process Look like?

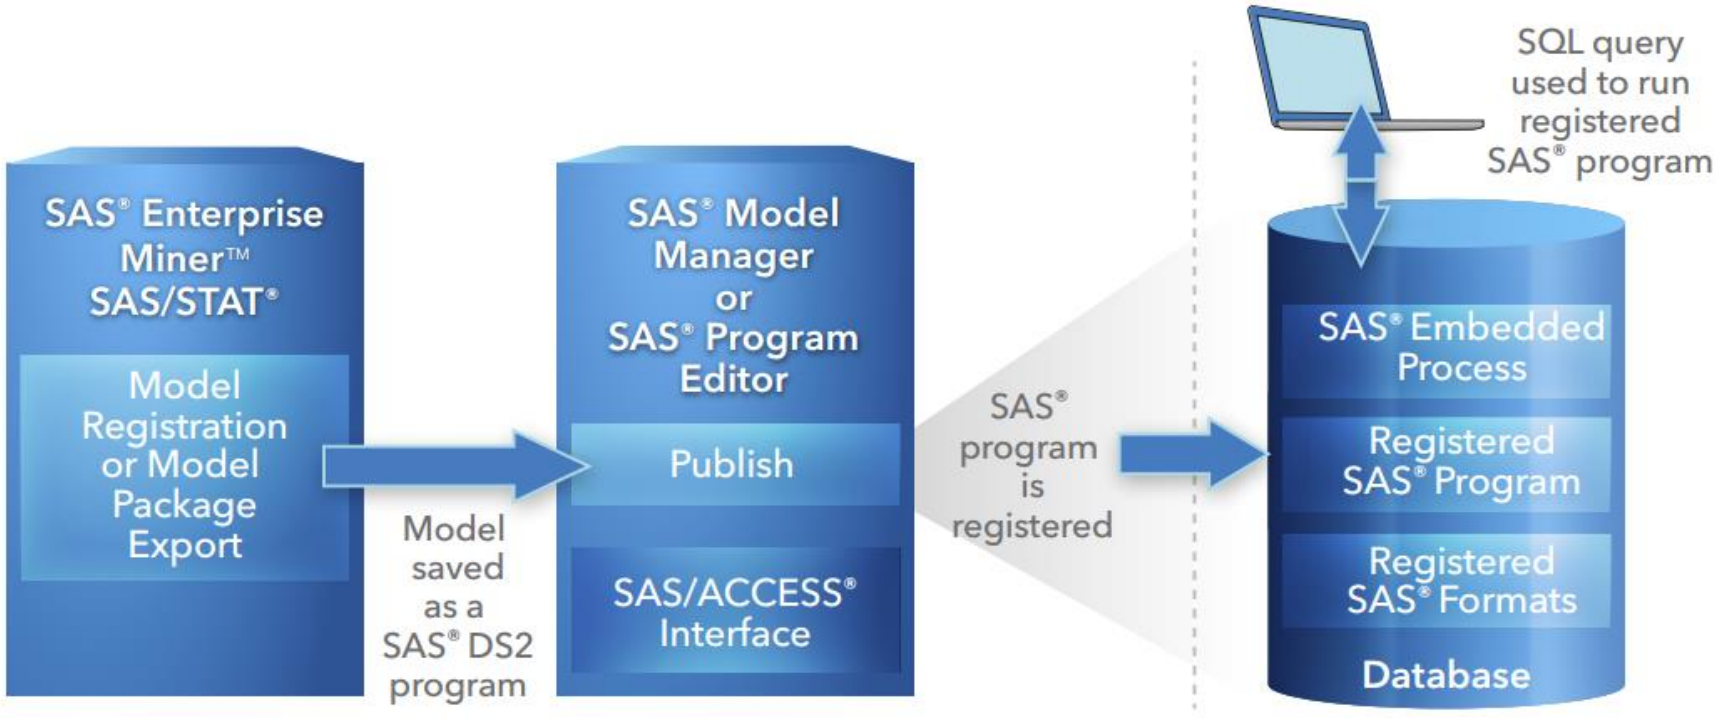

Data Loader for Hadoop

Data Loader for Hadoop – Self Service Big Data

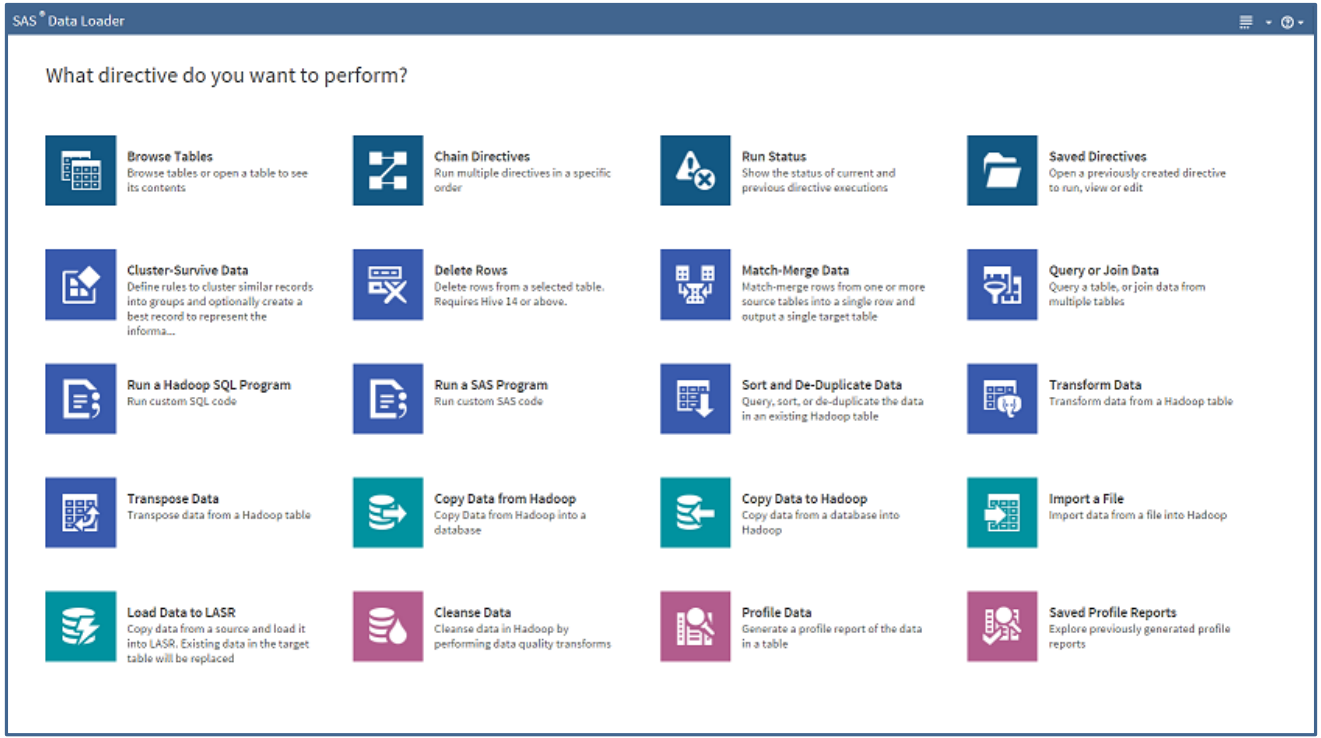

- Easy to use UI
- Query Data
- Manage Data
- Transform Data
- Run Custom Code
- Move Data

SAS Grid Manager for Hadoop

What is SAS Grid Manager for Hadoop?

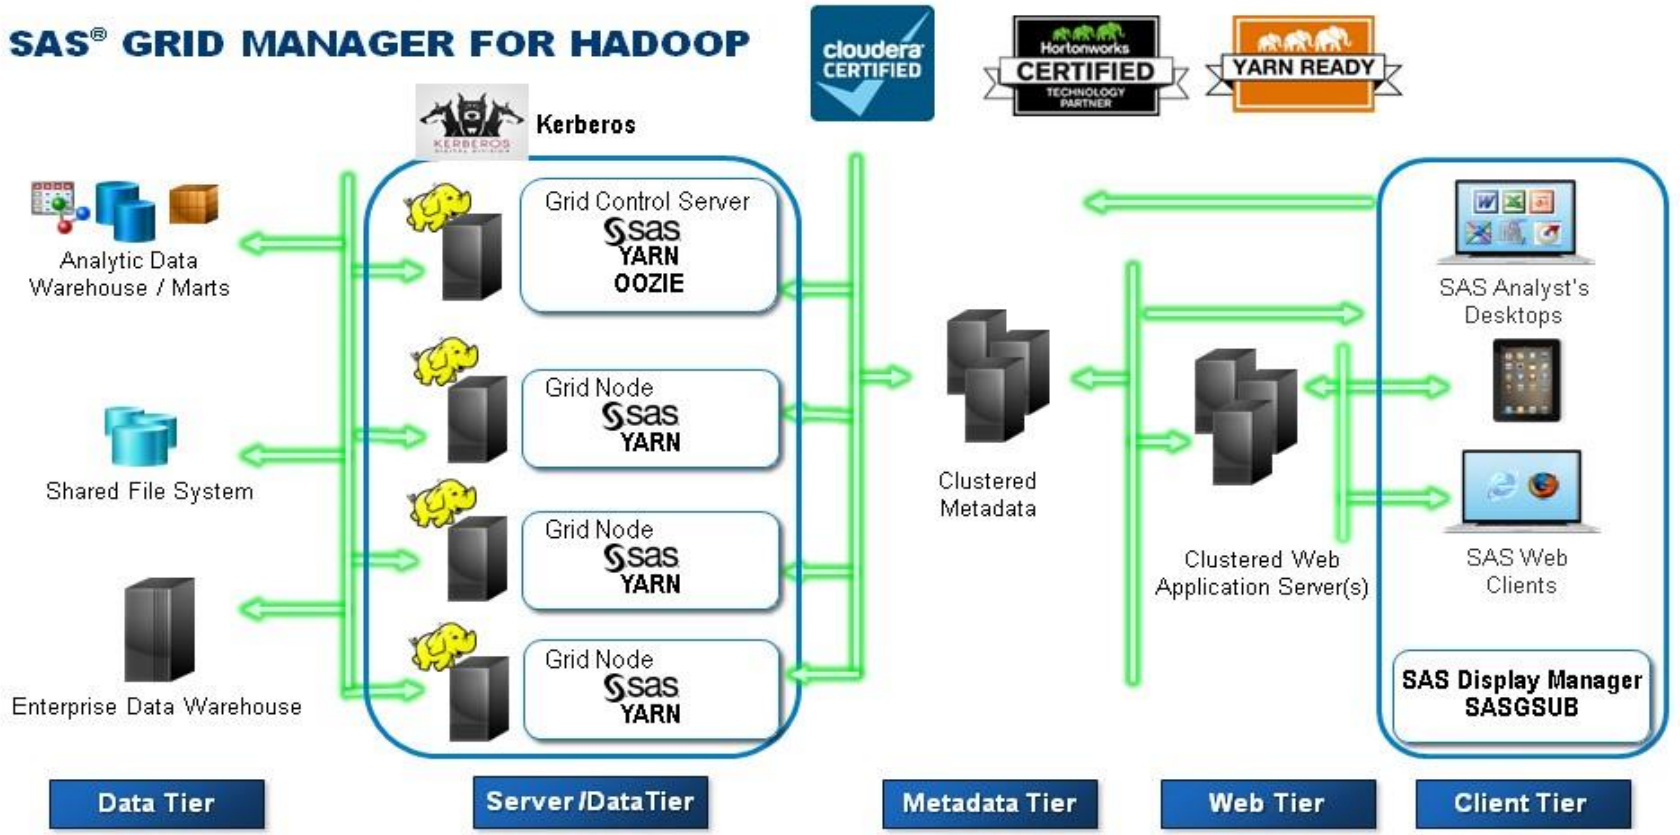

Servers

SAS Viya + Hadoop Architecture

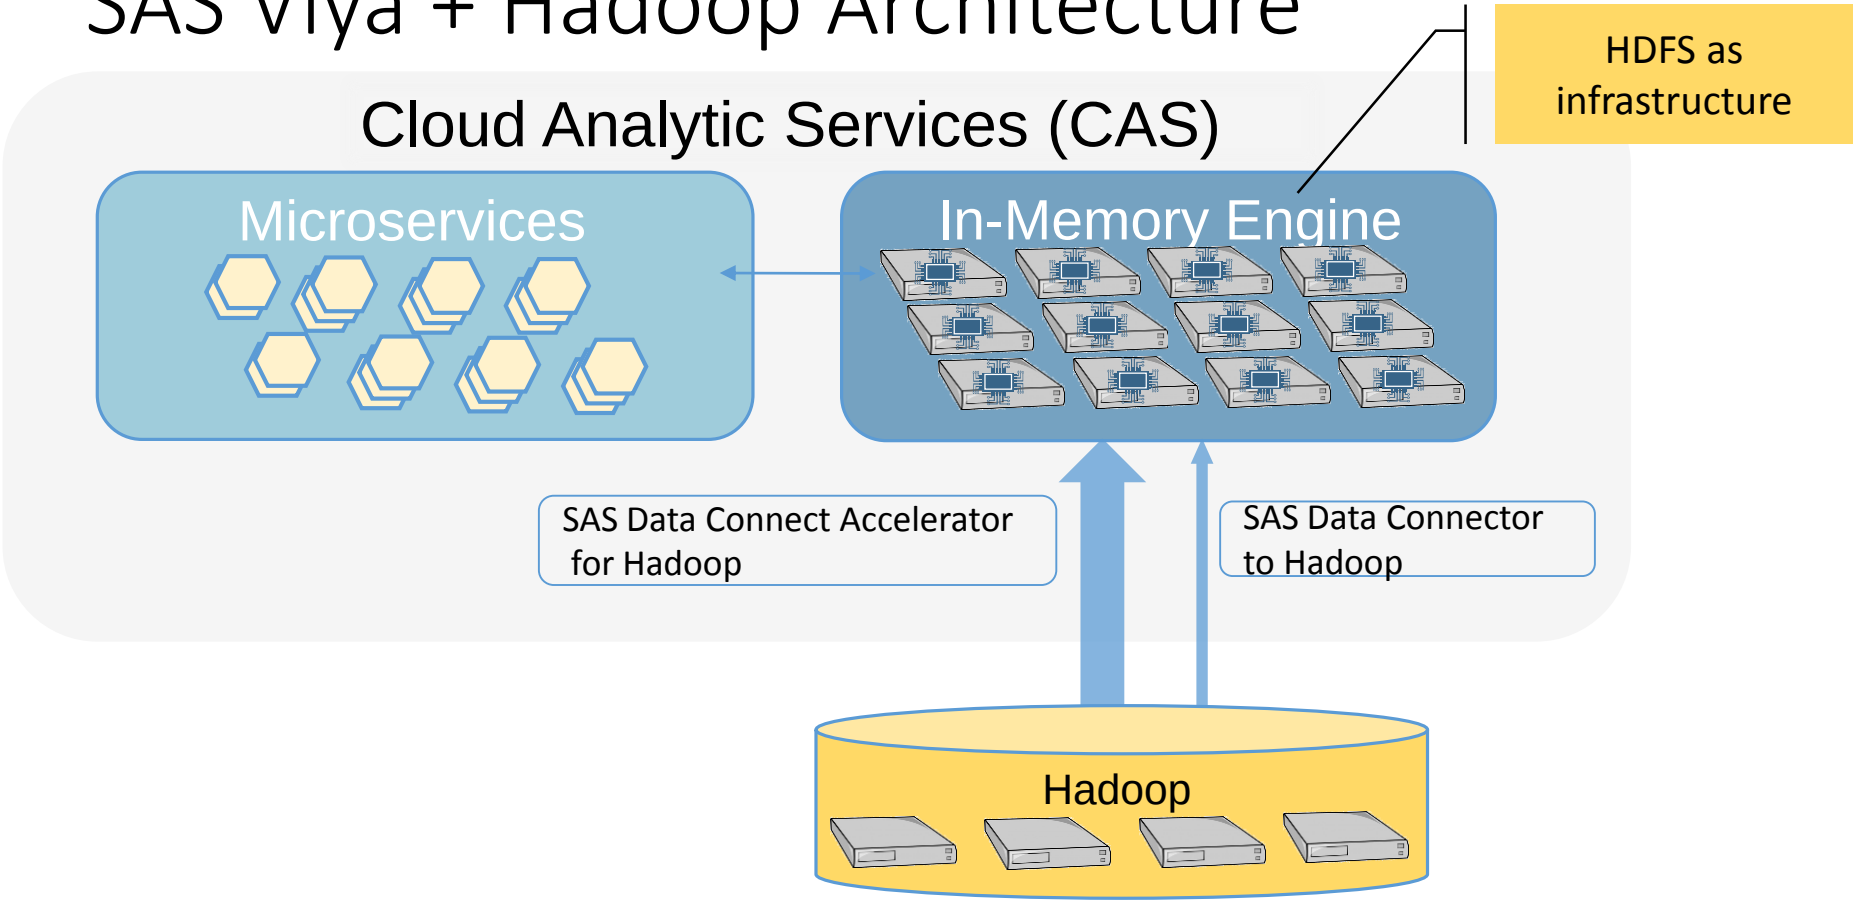

Feel Free to Contact Me!

Jeff.Bailey@sas.com

http://www.linkedin.com/in/jeffreydbailey

https://github.com/Jeff-Bailey/SAS13341_SAS_Hadoop

ANALYTICS EXPERIENCE 2016

#analyticsx